

Wishart Multivariate Response Linear Regression

John Tipton

December 13, 2015

Model

We begin with the model for the d -dimensional multivariate observation \mathbf{y}_i is

$$\mathbf{y}_i \sim \mathcal{N}(X_i \boldsymbol{\beta}, \boldsymbol{\Sigma}),$$

where $\boldsymbol{\beta}$ is the p -dimensional set of regression coefficients for covariate X_i and the covariance matrix is

$$\boldsymbol{\Sigma} = \begin{pmatrix} \sigma_1^2 & \rho_{1,2}\sigma_1\sigma_2 & \cdots & \rho_{1,d}\sigma_1\sigma_d \\ \rho_{1,2}\sigma_1\sigma_2 & \sigma_2^2 & \cdots & \rho_{2,d}\sigma_2\sigma_d \\ \vdots & \vdots & \ddots & \vdots \\ \rho_{1,d}\sigma_1\sigma_d & \rho_{2,d}\sigma_2\sigma_d & \cdots & \sigma_d^2 \end{pmatrix}.$$

We complete the model by assigning the priors

$$\begin{aligned} \boldsymbol{\beta} &\sim \mathcal{N}(\boldsymbol{\mu}_\beta, \boldsymbol{\Sigma}_\beta), \\ \rho_{j,k} &\sim \text{Unif}(-1, 1), \\ \sigma_j &\sim \text{Unif}(\sigma_l, \sigma_u), \end{aligned}$$

for $j, k = 1, \dots, d$ and $k > j$.

Posterior

The posterior that we wish to sample is

$$\begin{aligned}
[\boldsymbol{\beta}, \sigma_1, \dots, \sigma_d, \rho_{1,2}, \dots, \rho_{d-1,d} | \{\mathbf{y}_i, i = 1, \dots, N\}] &\propto \left(\prod_{i=1}^N [\mathbf{y}_i | \boldsymbol{\beta}, \sigma_1, \dots, \sigma_d, \rho_{1,2}, \dots, \rho_{d-1,d}] \right) [\boldsymbol{\beta}] \left(\prod_{j=1}^d [\sigma_j] \right) \\
&\times \left(\prod_{j=1}^d \prod_{k=j+1}^d [\rho_{j,k}] \right) \\
&\propto \left(\prod_{i=1}^N |\boldsymbol{\Sigma}|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (\mathbf{y}_i - X_i \boldsymbol{\beta})' \boldsymbol{\Sigma}^{-1} (\mathbf{y}_i - X_i \boldsymbol{\beta}) \right\} \right) \\
&\times |\boldsymbol{\Sigma}_{\boldsymbol{\beta}}|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (\boldsymbol{\beta} - \boldsymbol{\mu}_{\boldsymbol{\beta}})' \boldsymbol{\Sigma}_{\boldsymbol{\beta}}^{-1} (\boldsymbol{\beta} - \boldsymbol{\mu}_{\boldsymbol{\beta}}) \right\} \\
&\times \prod_{j=1}^d I \{ \sigma_l < \sigma_j < \sigma_u \} \\
&\times \prod_{j=1}^d \prod_{k=j+1}^d I \{ -1 \leq \rho_{j,k} \leq 1 \}
\end{aligned}$$

Full Conditionals

Full Conditional for $\boldsymbol{\beta}$

$$\begin{aligned}
[\boldsymbol{\beta} | \cdot] &\propto \prod_{i=1}^N [\mathbf{y}_i | \boldsymbol{\beta}, \boldsymbol{\Sigma}] [\boldsymbol{\beta}] \\
&\propto \prod_{i=1}^N \exp \left\{ -\frac{1}{2} (\mathbf{y}_i - \mathbf{X}_i \boldsymbol{\beta})' \boldsymbol{\Sigma}^{-1} (\mathbf{y}_i - \mathbf{X}_i \boldsymbol{\beta}) \right\} |\boldsymbol{\Sigma}_{\boldsymbol{\beta}}|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (\boldsymbol{\beta} - \boldsymbol{\mu}_{\boldsymbol{\beta}})' \boldsymbol{\Sigma}_{\boldsymbol{\beta}}^{-1} (\boldsymbol{\beta} - \boldsymbol{\mu}_{\boldsymbol{\beta}}) \right\} \\
&\propto \exp \left\{ -\frac{1}{2} \left(\boldsymbol{\beta}' \left(\sum_{i=1}^N \mathbf{X}_i' \boldsymbol{\Sigma}^{-1} \mathbf{X}_i + \boldsymbol{\Sigma}_{\boldsymbol{\beta}}^{-1} \right) \boldsymbol{\beta} - 2 \boldsymbol{\beta}' \left(\sum_{i=1}^N \mathbf{X}_i' \boldsymbol{\Sigma}^{-1} \mathbf{y}_i + \boldsymbol{\Sigma}_{\boldsymbol{\beta}}^{-1} \boldsymbol{\mu}_{\boldsymbol{\beta}} \right) \right) \right\}
\end{aligned}$$

which is $N(\mathbf{A}^{-1} \mathbf{b}, \mathbf{A}^{-1})$ with

$$\begin{aligned}
\mathbf{A} &= \sum_{i=1}^N \mathbf{X}_i' \boldsymbol{\Sigma}^{-1} \mathbf{X}_i + \boldsymbol{\Sigma}_{\boldsymbol{\beta}}^{-1} \\
\mathbf{b} &= \sum_{i=1}^N \mathbf{X}_i' \boldsymbol{\Sigma}^{-1} \mathbf{y}_i + \boldsymbol{\Sigma}_{\boldsymbol{\beta}}^{-1} \boldsymbol{\mu}_{\boldsymbol{\beta}}.
\end{aligned}$$

Full Conditional for σ_j

For $j = 1, \dots, d$,

$$[\sigma_j | \cdot] \propto \prod_{i=1}^N |\boldsymbol{\Sigma}|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (\mathbf{y}_i - \mathbf{X}_i \boldsymbol{\beta})' \boldsymbol{\Sigma}^{-1} (\mathbf{y}_i - \mathbf{X}_i \boldsymbol{\beta}) \right\} I \{ \sigma_l < \sigma_j < \sigma_u \}$$

which can be sampled using Metropolis-Hastings

Full Conditional for $\rho_{j,k}$

For $j, k = 1, \dots, d$ and $k > j$,

$$[\rho_{j,k}|\cdot] \propto \prod_{i=1}^N |\Sigma|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (\mathbf{y}_i - \mathbf{X}_i \beta)' \Sigma^{-1} (\mathbf{y}_i - \mathbf{X}_i \beta) \right\} I \{-1 \leq \rho_{j,k} \leq 1\}$$

which can be sampled using Metropolis-Hastings

Simulation and ‘R’ code

```
set.seed(101)
N <- 100
d <- 2
beta <- rnorm(d)
X <- array(rnorm(N * d^2), dim = c(N, d, d))
sigma_l <- 0
sigma_u <- 10
sigma <- runif(d, sigma_l, sigma_u)
rho <- runif(d * (d - 1)/2, -1, 1)

## Construct a generic dxd covariance matrix
makeCov <- function(d, sigma, rho) {
  sigmaMat <- sigma %*% t(sigma)
  rhoMat <- diag(d)
  rhoMat[lower.tri(rhoMat)] <- rho
  rhoMat <- rhoMat + t(rhoMat) - diag(d)
  # rhoMat[upper.tri(rhoMat)] <- t(rhoMat[lower.tri(rhoMat)])
  covMat <- sigmaMat * rhoMat
  return(covMat)
}

Sigma <- makeCov(d, sigma, rho)
## Simulate data
library(mvtnorm)
y <- matrix(0, N, d)
for (i in 1:N) {
  y[i, ] <- rmvnorm(1, X[i, , ] %*% beta, Sigma)
}

## MCMC function in R
mcmcR <- function(n_mcmc, y, mu_beta, Sigma_beta, sigma_l, sigma_u, sigma_tune = 1,
  rho_tune = 0.1) {

  library(mvtnorm)
  ## set up dimensions
```

```

N <- dim(y)[1]
d <- dim(y)[2]

## setup save variables
beta = matrix(0, n_mcmc, d)
sigma <- matrix(0, n_mcmc, d)
rho <- matrix(0, n_mcmc, d * (d - 1)/2)
Sigma_inv <- array(0, dim = c(n_mcmc, d, d))

## initialize values
beta[1, ] <- rmvnorm(1, mu_beta, Sigma_beta)
rho[1, ] <- runif(d * (d - 1)/2, -1, 1)
sigma[1, ] <- runif(d, sigma_l, sigma_u)
Sigma_inv[1, , ] <- solve(makeCov(d, sigma[1, ], rho[1, ]))
Sigma_beta_inv <- solve(Sigma_beta)
y_sum <- apply(y, 2, sum)
ty <- t(y)
sigma_tune <- rep(sigma_tune, d)
sigma_accept_tmp <- rep(0, d)
sigma_accept <- rep(0, d)
rho_tune <- rep(rho_tune, d * (d - 1)/2)
rho_accept_tmp <- rep(0, d * (d - 1)/2)
rho_accept <- rep(0, d * (d - 1)/2)

message(paste("Starting MCMC fit, will run for", n_mcmc, "iterations"))

## Start MCMC chain
for (k in 2:n_mcmc) {
  if (k%%500 == 0) {
    message(paste("Iteration", k))
  }

  ## sample beta
  A_inv <- solve(Reduce("+", lapply(seq_len(dim(X)[1]), function(i) {
    t(X[i, , ]) %*% Sigma_inv[k - 1, , ] %*% X[i, , ]
  }))) + Sigma_beta_inv)
  b <- Reduce("+", lapply(seq_len(dim(X)[1]), function(i) {
    t(X[i, , ]) %*% Sigma_inv[k - 1, , ] %*% y[i, ]
  }))) + Sigma_beta_inv %*% mu_beta
  beta[k, ] <- rmvnorm(1, A_inv %*% b, A_inv)

  ## sample sigma
  Sigma_inv[k, , ] <- Sigma_inv[k - 1, , ]
  sigma[k, ] <- sigma[k - 1, ]
  for (j in 1:d) {
    sigma_star <- sigma[k, ]
    sigma_star[j] <- rnorm(1, sigma[k, j], sigma_tune[j])
    if (sigma_star[j] > sigma_l && sigma_star[j] < sigma_u) {
      Sigma_inv_star <- solve(makeCov(d, sigma_star, rho[k - 1, ]))
      mh1 <- N * sum(log(diag(chol(Sigma_inv_star)))) - 0.5 * sum(unlist(lapply(seq_len(dim(X)),
        function(i) {
          t(y[i, ] - X[i, , ] %*% beta[k, ]) %*% Sigma_inv_star %*%

```

```

      (y[i, ] - X[i, , ] %*% beta[k, ])
    })))
  mh2 <- N * sum(log(diag(chol(Sigma_inv[k, , ])))) - 0.5 * sum(unlist(lapply(seq_len(dim(
    function(i) {
      t(y[i, ] - X[i, , ] %*% beta[k, ]) %*% Sigma_inv[k, , ] %*%
        (y[i, ] - X[i, , ] %*% beta[k, ])
    })))
  mh <- exp(mh1 - mh2)
  if (mh > runif(1)) {
    sigma[k, ] <- sigma_star
    Sigma_inv[k, , ] <- Sigma_inv_star
    sigma_accept_tmp[j] <- sigma_accept_tmp[j] + 1/50
    sigma_accept[j] <- sigma_accept[j] + 1/n_mcmc
  }
}

## Update tuning
if (k%%50 == 1) {
  for (j in 1:d) {
    if (sigma_accept_tmp[j] > 0.44) {
      sigma_tune[j] <- exp(log(sigma_tune[j]) + 1/sqrt(k))
    } else {
      sigma_tune[j] <- exp(log(sigma_tune[j]) - 1/sqrt(k))
    }
    sigma_accept_tmp[j] <- 0
  }
}

## sample rho
rho[k, ] <- rho[k - 1, ]
for (j in 1:(d * (d - 1)/2)) {
  rho_star <- rho[k, ]
  rho_star[j] <- rnorm(1, rho[k, j], rho_tune[j])
  if (rho_star[j] >= -1 && rho_star[j] <= 1) {
    Sigma_inv_star <- solve(makeCov(d, sigma[k, ], rho_star))
    mh1 <- N * sum(log(diag(chol(Sigma_inv_star)))) - 0.5 * sum(unlist(lapply(seq_len(dim(X
      function(i) {
        t(y[i, ] - X[i, , ] %*% beta[k, ]) %*% Sigma_inv_star %*%
          (y[i, ] - X[i, , ] %*% beta[k, ])
      })))
    mh2 <- N * sum(log(diag(chol(Sigma_inv[k, , ])))) - 0.5 * sum(unlist(lapply(seq_len(dim(
      function(i) {
        t(y[i, ] - X[i, , ] %*% beta[k, ]) %*% Sigma_inv[k, , ] %*%
          (y[i, ] - X[i, , ] %*% beta[k, ])
      })))
    mh <- exp(mh1 - mh2)
    if (mh > runif(1)) {
      rho[k, ] <- rho_star
      Sigma_inv[k, , ] <- Sigma_inv_star
      rho_accept_tmp[j] <- rho_accept_tmp[j] + 1/50
      rho_accept[j] <- rho_accept[j] + 1/n_mcmc
    }
  }
}

```

```

    }
  }
  ## Update tuning
  if (k%%50 == 1) {
    for (j in 1:(d * (d - 1)/2)) {
      if (rho_accept_tmp[j] > 0.44) {
        rho_tune[j] <- exp(log(rho_tune[j]) + 1/sqrt(k))
      } else {
        rho_tune[j] <- exp(log(rho_tune[j]) - 1/sqrt(k))
      }
      rho_accept_tmp[j] <- 0
    }
  }
}

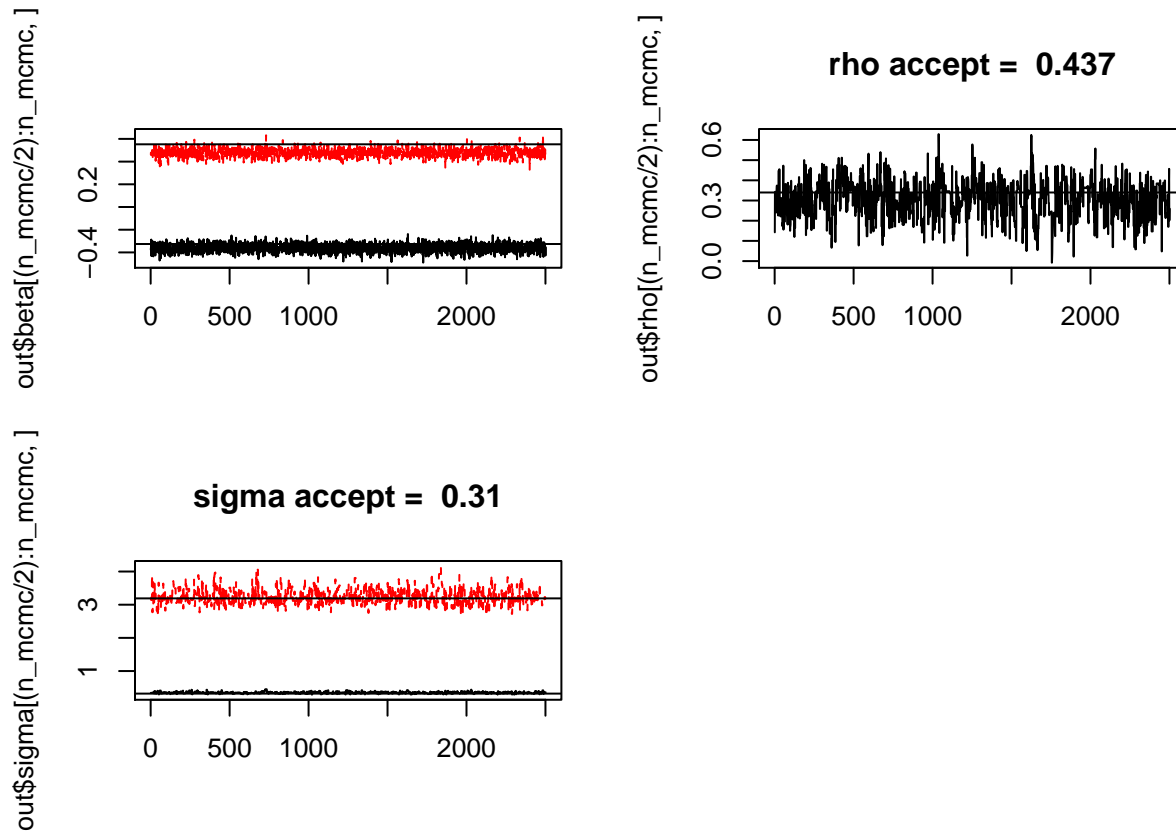
## Output MCMC
return(list(beta = beta, sigma = sigma, rho = rho, Sigma_inv = Sigma_inv,
           sigma_accept = sigma_accept, rho_accept = rho_accept))
}

## Define priors
mu_beta <- rep(0, d)
Sigma_beta <- 100 * diag(d)
sigma_l <- 0
sigma_u <- 10
n_mcmc <- 5000

## Run MCMC
out <- mcmcR(n_mcmc, y, mu_beta, Sigma_beta, sigma_l, sigma_u)

## Plot MCMC output (post burn-in)
layout(matrix(1:4, 2, 2))
matplot(out$beta[(n_mcmc/2):n_mcmc, ], type = 'l')
abline(h=beta)
matplot(out$sigma[(n_mcmc/2):n_mcmc, ], type = 'l', main=paste("sigma accept = ", round(mean(out$sigma_
abline(h=sigma)
matplot(out$rho[(n_mcmc/2):n_mcmc, ], type = 'l', main=paste("rho accept = ", round(mean(out$rho_accept
abline(h=rho)

```



```
## Compare Estimates (post burn-in) to truth
beta
```

```
## [1] -0.3260365  0.5524619
```

```
apply(out$beta[(n_mcmc/2):n_mcmc, ], 2, mean)
```

```
## [1] -0.3623669  0.4783602
```

```
sigma
```

```
## [1] 0.3202498 3.1906751
```

```
apply(out$sigma[(n_mcmc/2):n_mcmc, ], 2, mean)
```

```
## [1] 0.3518111 3.2408612
```

```
rho
```

```
## [1] 0.3398204
```

```
if(dim(out$rho)[2] == 1){  
  mean(out$rho[(n_mcmc/2):n_mcmc, ])  
} else {  
  apply(out$rho[(n_mcmc/2):n_mcmc, ], 2, mean)  
}
```

```
## [1] 0.3059642
```